

# Linux Device Drivers

## On Desktops

Team Emertxe



# Introduction

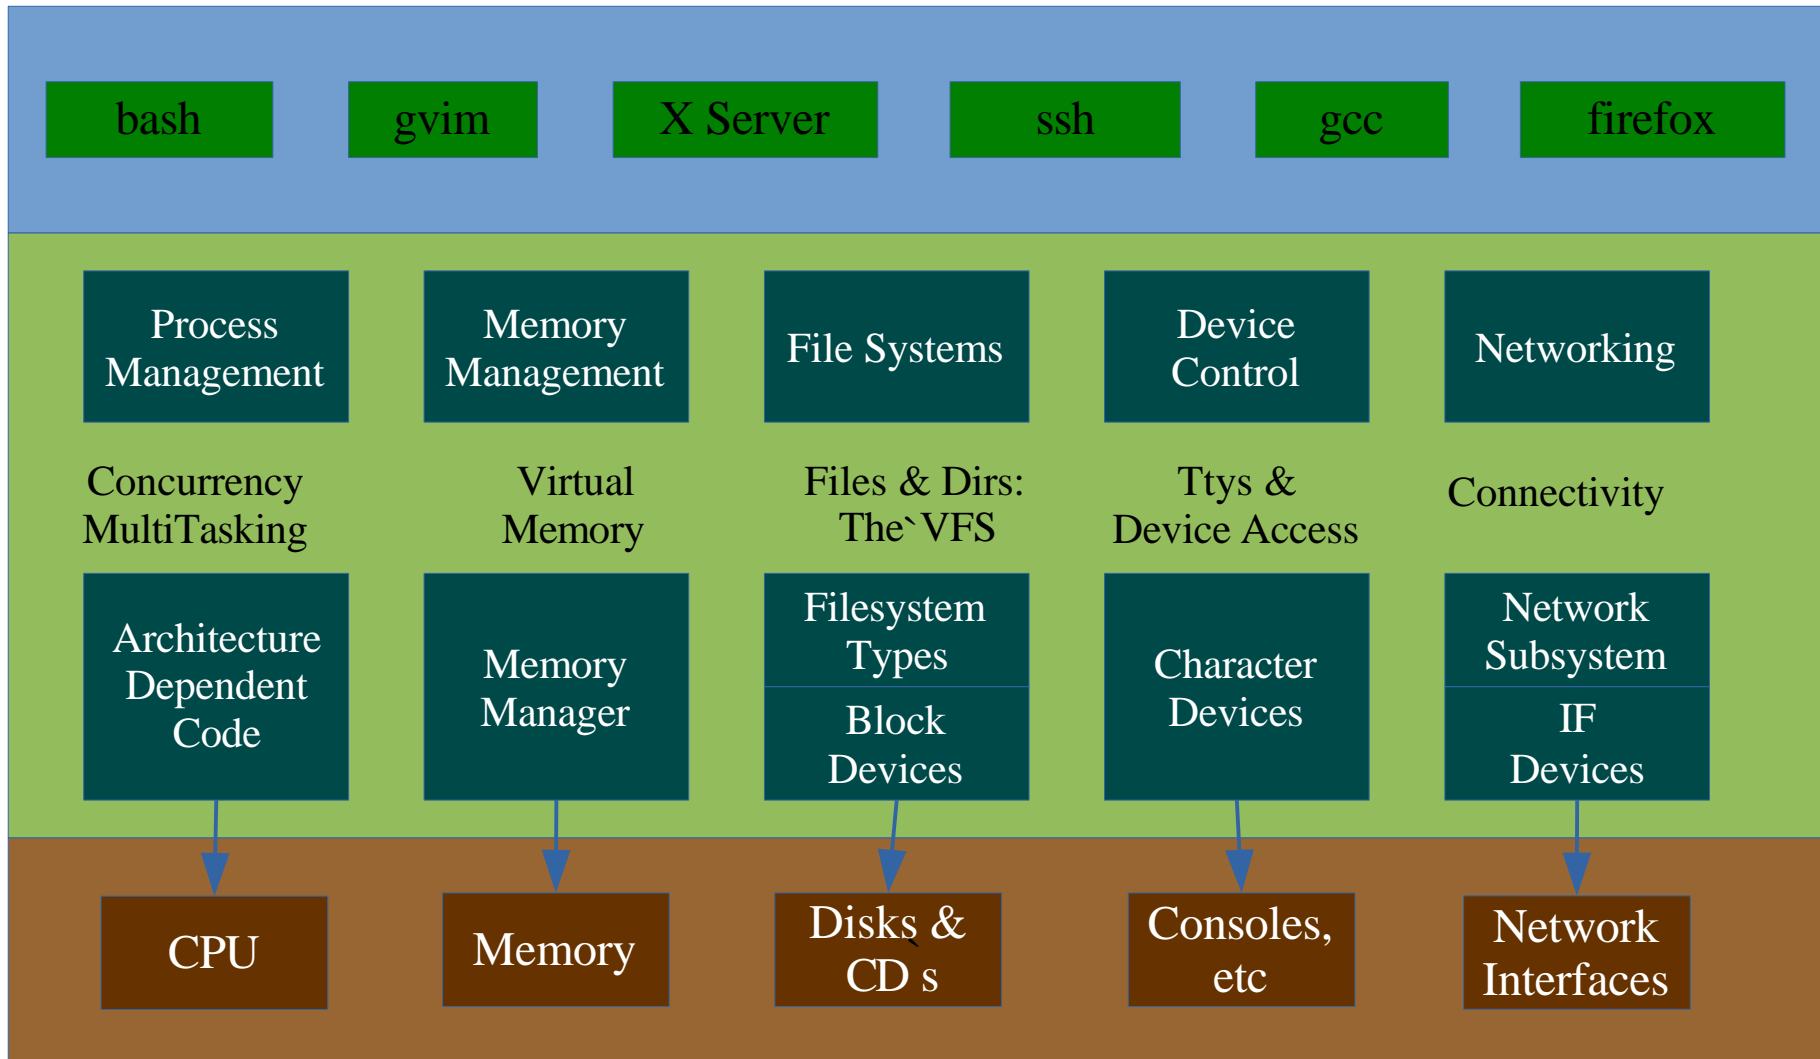


# Familiarity Check



- Good C & Programming Skills
- Linux & the Filesystem
  - Root, User Space Headers & Libraries
- Files
  - Regular, Special, Device
- Toolchain
  - gcc & friends
- Make & Makefiles
- Kernel Sources (Location & Building)

# Linux Driver Ecosystem





# The Flow



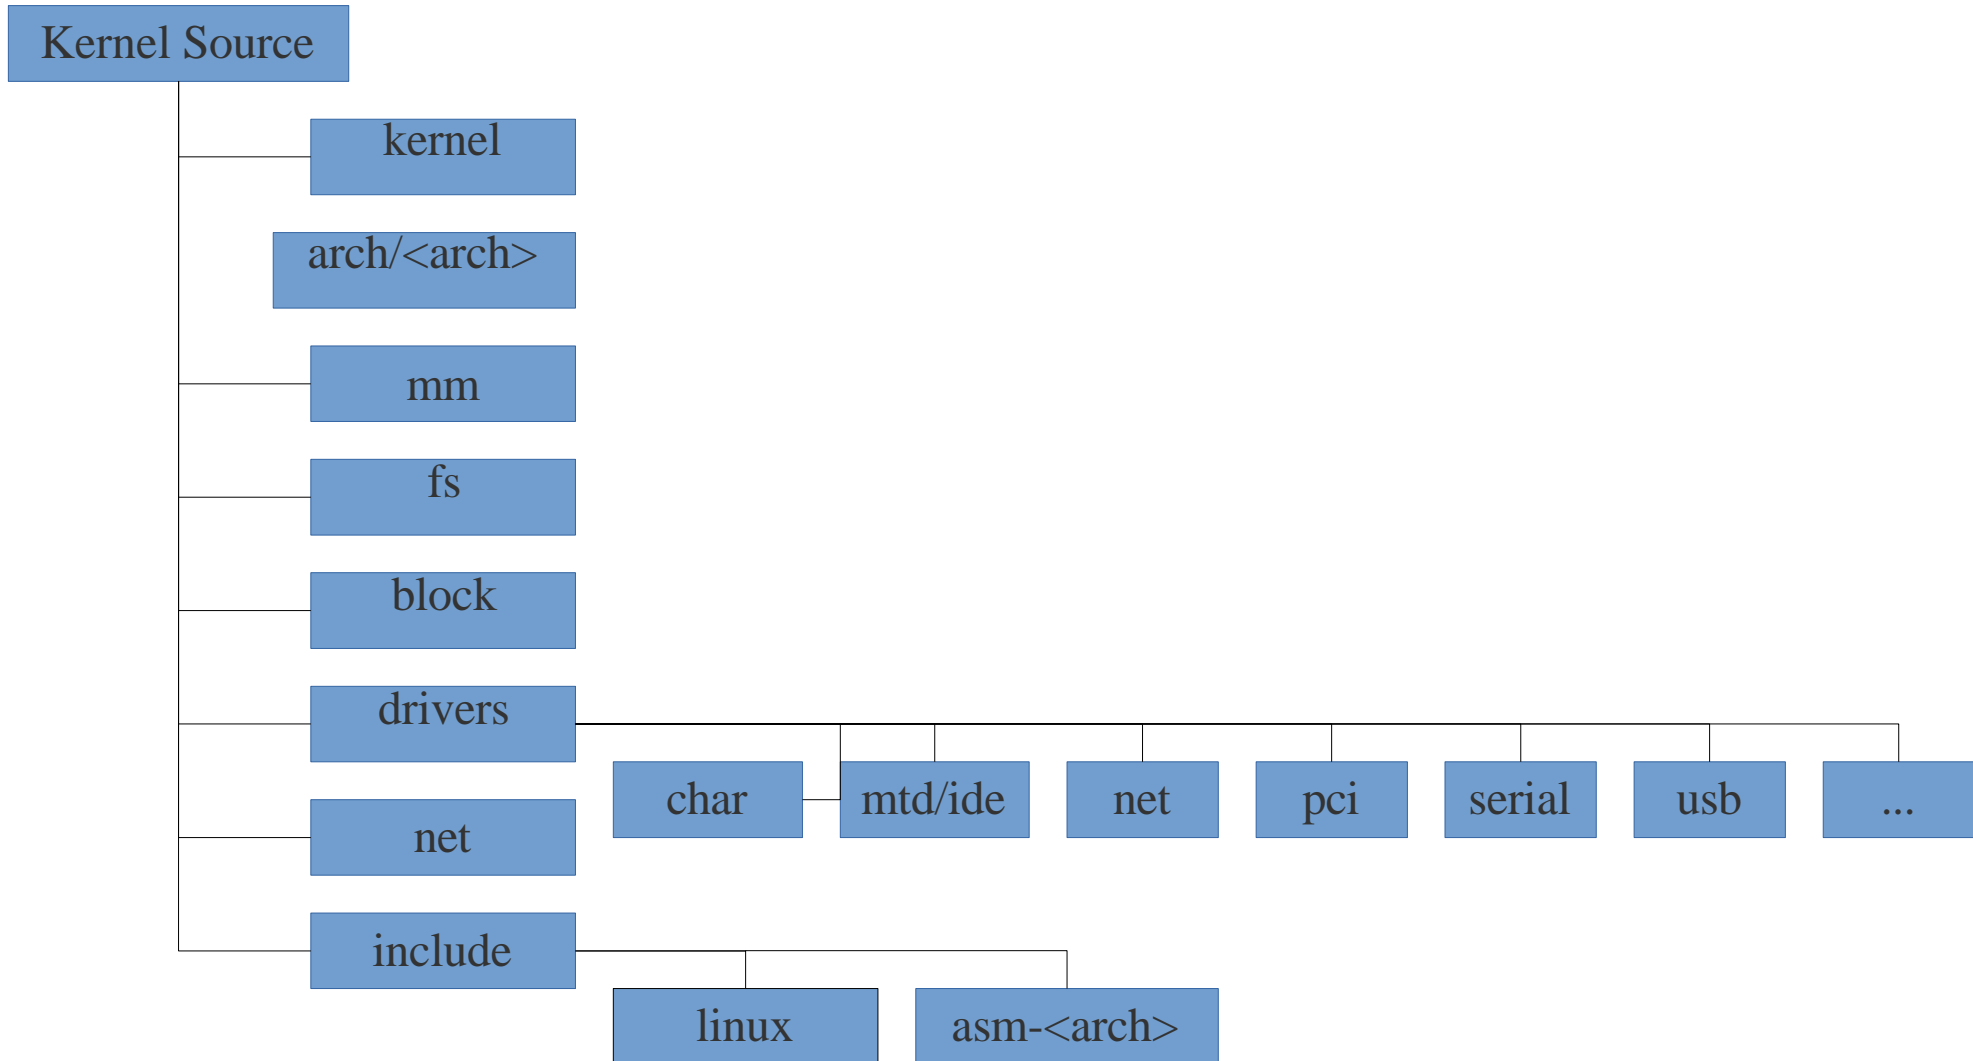
- Introduction
- Character Drivers
- Memory & Hardware
- Time & Timings
- USB Drivers
- Interrupt Handling
- Block Drivers
- PCI Drivers

# Hands-On



- Your First Driver
- Character Drivers
  - Null Driver
  - Memory Driver
  - UART Driver for Customized Hardware
- USB Drivers
  - USB Device Hot-plug-ability
  - USB to Serial Hardware Driver
- Filesystem Modules
  - VFS Interfacing
  - “Pseudo” File System with Memory Files

# Kernel Source Organization



# The Locations & Config Files



- Kernel Source Path: `/usr/src/linux`
- Std Modules Path:
  - `/lib/modules/<kernel version>/kernel/...`
- Module Configuration: `/etc/modprobe.conf`
- Kernel Windows:
  - `/proc`
  - `/sys`
- System Logs: `/var/log/messages`





# The Commands



- lsmod
- insmod
- modprobe
- rmmod
- dmesg
- objdump
- nm
- cat /proc/<file>

# The Kernel's C

- ctor & dtor
  - init\_module, cleanup\_module
- printf
  - printk
- Libraries
  - <kernel src>/kernel
- Headers
  - <kernel src>/include

# The Init Code

```
static int __init mfd_init(void)
{
    printk(KERN_INFO "mfd
        registered");
    ...
    return 0;
}
module_init(mfd_init);
```

# The Cleanup Code

```
static void __exit mfd_exit(void)
{
    printk(KERN_INFO "mfd
        deregistered");
    ...
}
module_exit(mfd_exit);
```

# Usage of printk

- `<linux/kernel.h>`
- Constant String for Log Level
  - `KERN_EMERG`      "`<0>`"    `/* system is unusable */`
  - `KERN_ALERT`      "`<1>`"    `/* action must be taken immediately */`
  - `KERN_CRIT`        "`<2>`"    `/* critical conditions */`
  - `KERN_ERR`         "`<3>`"    `/* error conditions */`
  - `KERN_WARNING`    "`<4>`"    `/* warning conditions */`
  - `KERN_NOTICE`     "`<5>`"    `/* normal but significant condition */`
  - `KERN_INFO`        "`<6>`"    `/* informational */`
  - `KERN_DEBUG`      "`<7>`"    `/* debug-level messages */`
- `printf` like arguments

# The Other Basics & Ornaments



- Headers
  - `#include <linux/module.h>`
  - `#include <linux/version.h>`
  - `#include <linux/kernel.h>`
- `MODULE_LICENSE("GPL");`
- `MODULE_AUTHOR("Emertxe");`
- `MODULE_DESCRIPTION("First Device Driver");`

# Building the Module



- Our driver needs
  - The Kernel Headers for Prototypes
  - The Kernel Functions for Functionality
  - The Kernel Build System & the Makefile for Building
- Two options
  - Building under Kernel Source Tree
    - Put our driver under drivers folder
    - Edit Kconfig(s) & Makefile to include our driver
  - Create our own Makefile to do the right invocation

# Our Makefile

```
ifneq (${KERNELRELEASE},)
    objm += <module>.o
else
    KERNEL_SOURCE := <kernel source directory path>
    PWD := $(shell pwd)
default:
    $(MAKE) C ${KERNEL_SOURCE} M=$(PWD)
    modules
clean:
    $(MAKE) C ${KERNEL_SOURCE} M=$(PWD)
    clean
endif
```



Try Out your First Driver



# Character Drivers



# Major & Minor Number



- `ls -l /dev`
- Major is to Driver; Minor is to Device
- `<linux/types.h>` (`>= 2.6.0`)
  - `dev_t`: 12 & 20 bits for major & minor
- `<linux/kdev_t.h>`
  - `MAJOR(dev_t dev)`
  - `MINOR(dev_t dev)`
  - `MKDEV(int major, int minor)`

# Registering & Unregistering



- Registering the Device Driver
  - `int register_chrdev_region(dev_t first, unsigned int count, char *name);`
  - `int alloc_chrdev_region(dev_t *dev, unsigned int firstminor, unsigned int cnt, char *name);`
- Unregistering the Device Driver
  - `void unregister_chrdev_region(dev_t first, unsigned int count);`
- Header: `<linux/fs.h>`



# The file operations

- `#include <linux/fs.h>`
- `struct file_operations`
  - `int (*open)(struct inode *, struct file *);`
  - `int (*release)(struct inode *, struct file *);`
  - `ssize_t (*read)(struct file *, char __user *, size_t, loff_t *);`
  - `ssize_t (*write)(struct file *, const char __user *, size_t, loff_t *);`
  - `struct module owner = THIS_MODULE; / linux/module.h> */`
  - `loff_t (*llseek)(struct file *, loff_t, int);`
  - `int (*ioctl)(struct inode *, struct file *, unsigned int, unsigned long);`

# User level I/O

- `int open(const char *path, int oflag, ... )`
- `int close(int fd);`
- `ssize_t write(int fd, const void *buf, size_t nbyte)`
- `ssize_t read(int fd, void *buf, size_t nbyte)`
- `int ioctl(int d, int request, ...)`
  - The `ioctl()` function manipulates the underlying device parameters of special files.
  - The argument `d` must be an open file descriptor.
  - The second argument is a device-dependent request code.

# The file & inode structures



- struct file
  - mode\_t f\_mode
  - loff\_t f\_pos
  - unsigned int f\_flags
  - struct file\_operations \*f\_op
  - void \* private\_data
- struct inode
  - unsigned int iminor(struct inode \*);
  - unsigned int imajor(struct inode \*);

# Registering the file operations



- `#include <linux/cdev.h>`
- 1<sup>st</sup> way initialization:
  - `struct cdev *my_cdev = cdev_alloc();`
  - `my_cdev->owner = THIS_MODULE;`
  - `my_cdev->ops = &my_fops;`
- 2<sup>nd</sup> way initialization:
  - `struct cdev my_cdev;`
  - `cdev_init(&my_cdev, &my_fops);`
  - `my_cdev.owner = THIS_MODULE;`
  - `my_cdev.ops = &my_fops;`



# Registering the file operations...



- The Registration
  - `int cdev_add(struct cdev *cdev, dev_t num, unsigned int count);`
- The Unregistration
  - `void cdev_del(struct cdev *cdev);`



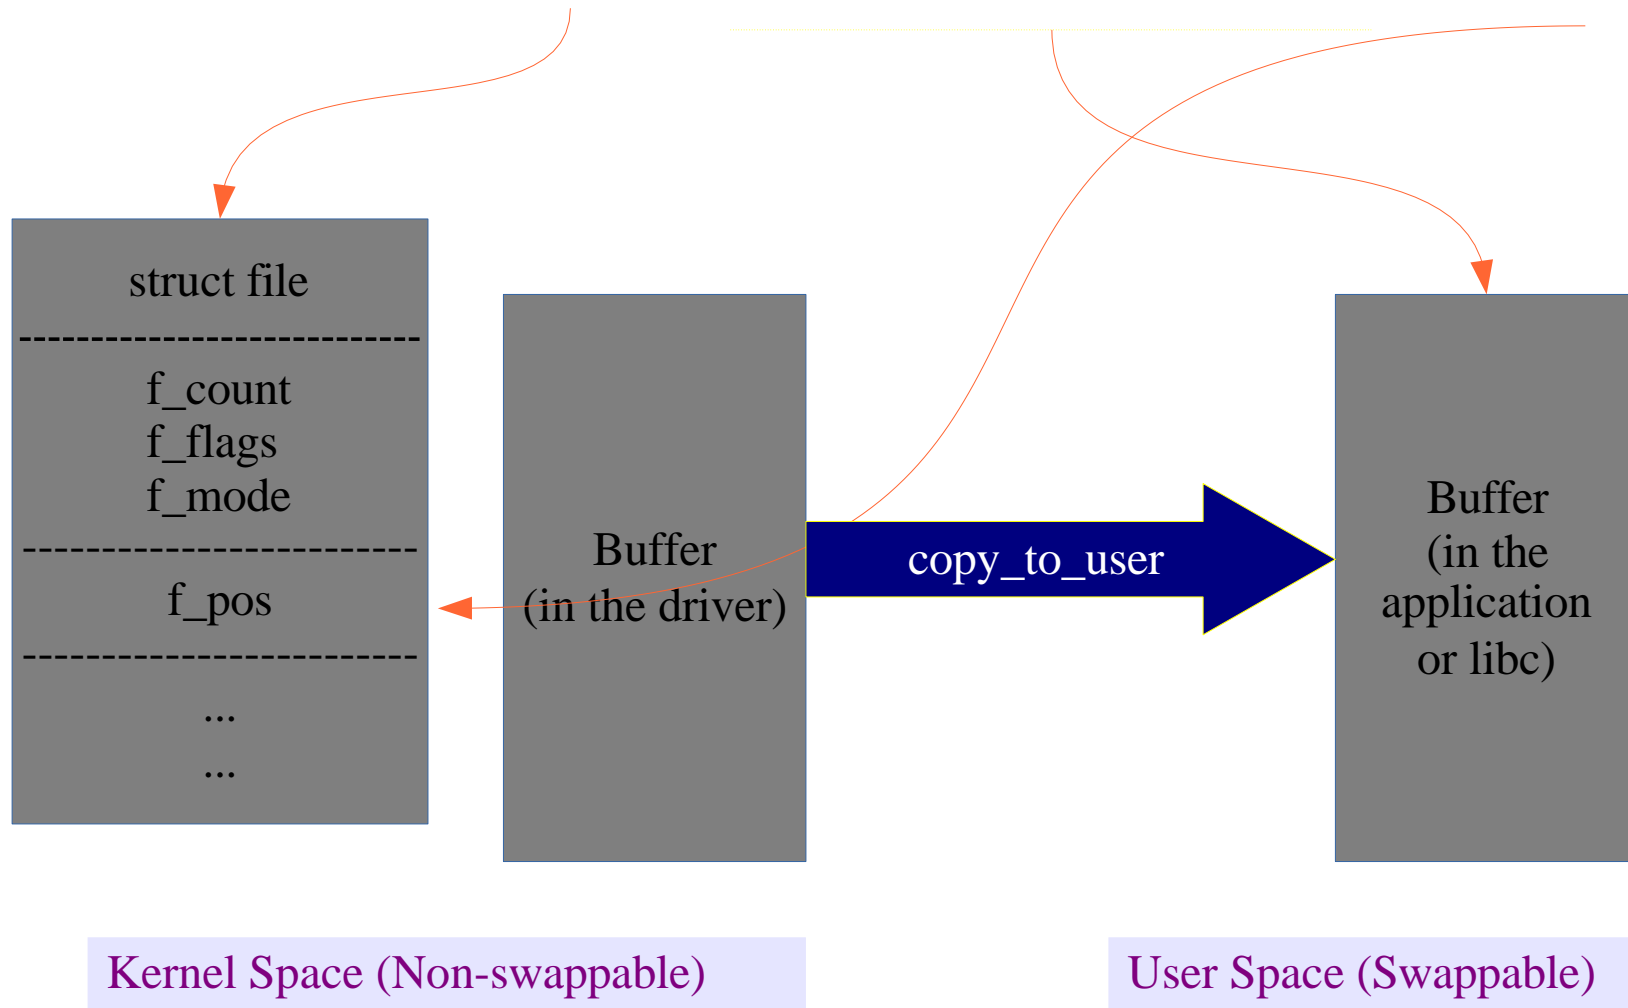
# Registering/Unregistering Old Way



- Registering the Device Driver
  - `int register_chrdev(undigned int major, const char *name, struct file_operations *fops);`
- Unregistering the Device Driver
  - `int unregister_chrdev(undigned int major, const char *name);`

# The read flow

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t *off)
```



# The /dev/null read & write



```
ssize_t my_read(struct file *f, char __user
    *buf, size_t cnt, loff_t *off)
{
    ...
    return read_cnt;
}
ssize_t my_write(struct file *f, char __user
    *buf, size_t cnt, loff_t *off)
{
    ...
    return wrote_cnt;
}
```

# The mem device read

```
ssize_t my_read(struct file *f, char __user
    *buf, size_t cnt, loff_t *off)
{
    ...
    if (copy_to_user(buf, from, cnt) != 0)
    {
        return EFAULT;
    }
    ...
    return read_cnt;
}
```

# The mem device write

```
ssize_t my_write(struct file *f, char __user
    *buf, size_t cnt, loff_t *off)
{
    ...
    if (copy_from_user(to, buf, cnt) != 0)
    {
        return EFAULT;
    }
    ...
    return wrote_cnt;
}
```

# Dynamic Device Node & Classes



- Class Operations

- `struct class *class_create(struct module *owner, char *name);`
- `void class_destroy(struct class *cl);`

- Device into & Out of Class

- `struct class_device *device_create(struct class *cl, NULL, dev_t devnum, NULL, const char *fmt, ...);`
- `void device_destroy(struct class *cl, dev_t devnum);`

# The I/O Control API



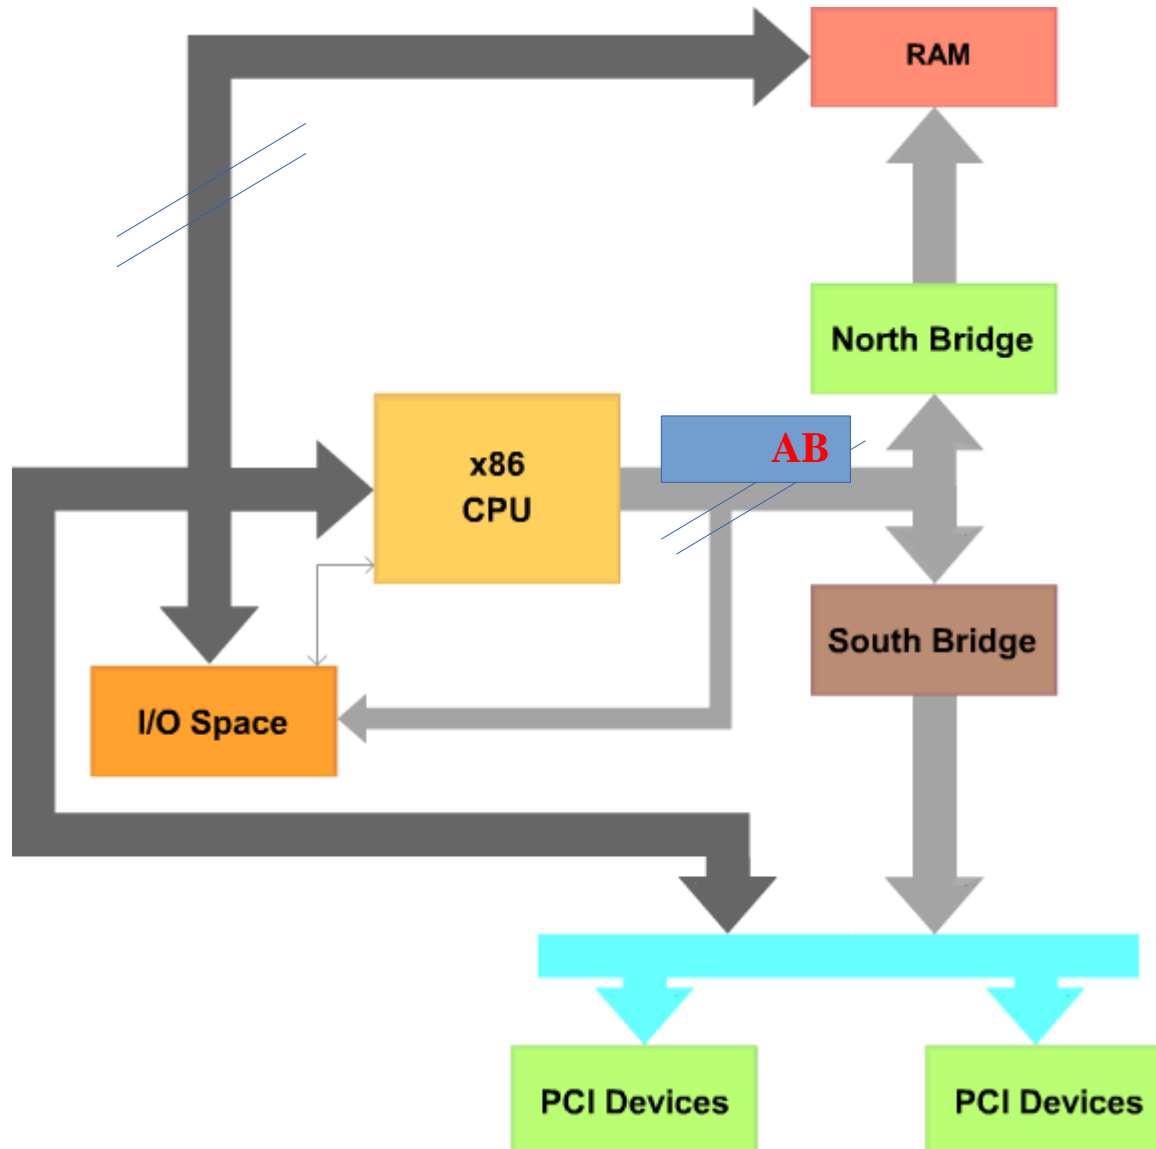
- `int (*ioctl)(struct inode *, struct file *, unsigned int cmd, unsigned long arg)`
- Command
  - `<linux/ioctl.h> -> ... -> <asm-generic/ioctl.h>`
  - Macros
    - `_IO, _IOR, _IOW, _IOWR`
  - Parameters
    - type (character) [15:8]
    - number (index) [7:0]
    - size (param type) [29:16]



# Module Parameters

- <linux/moduleparam.h>
  - Macros
    - module\_param(name, type, perm)
    - module\_param\_array(name, type, num, perm)
    - Perm (is a bitmask)
      - 0
      - S\_IRUGO
      - S\_IWUSR | S\_IRUGO
  - Loading
    - insmod driver.ko name=10

# x86 Architecture




# Memory Access



# Physical Vs Virtual Memory



- The kernel Organizes Physical memory in to pages
  - Page size Depends on Arch
    - X86-based 4096 bytes
- On 32-bit X86 system Kernel total Virtual address space
  - Total 4GB (pointer size)
  - Kernel Configuration Splits 4GB in to
    - 3GB Virtual Sp for US
    - 1GB Virtual Sp for Kernel
      - 128MB KDS
  - Virtual Address also called Address”



# Memory Access from Kernel Space



- Virtual Address on Physical Address
  - `#include <linux/gfp.h>`
    - `unsigned long __get_free_pages(flags, order);` etc
    - `void free_pages(addr, order);` etc
  - `#include <linux/slab.h>`
    - `void *kmalloc(size_t size, gfp_t flags);`
      - `GFP_ATOMIC`, `GFP_KERNEL`, `GFP_DMA`
    - `void kfree(void *obj);`
  - `#include <linux/vmalloc.h>`
    - `void *vmalloc(unsigned long size);`
    - `void vfree(void *addr);`

# Memory Access from Kernel Space...



- Virtual Address for Bus/IO Address
  - `#include <asm/io.h>`
    - `void *ioremap(unsigned long offset, unsigned long size);`
    - `void iounmap(void *addr);`
- I/O Memory Access
  - `#include <asm/io.h>`
    - `unsigned int ioread[8|16|32](void *addr);`
    - `unsigned int iowrite[8|16|32](u[8|16|32] value, void *addr);`
- Barriers
  - `#include <linux/kernel.h>: void barrier(void);`
  - `#include <asm/system.h>: void [r|w|]mb(void);`

# Debugging



# Debugging Options

- Printing & syslogd
- Querying
  - /proc
  - /sysfs

## Watching

- - strace
  - Oops



# Introduction to Procfs



## Objective:

- Understand the purpose of procfs in Linux.
- Learn how to create read/write interfaces using procfs.
- Explore practical examples for device drivers.

# What is Procfs?

- **Definition:**

- Procfs is a virtual file system in Linux that provides
- information about processes and kernel data.

- **Mounted Location:**

- Usually mounted at /proc.

- **Purpose:**

- Exposes process and system information.
- Provides an interface for user-space programs
- to interact with kernel data.

# Why Use Procfs?

- **Advantages:**

- Provides system information in a human-readable format.
- Allows users to configure kernel parameters.

- **Common Use Cases:**

- Viewing process information.
- Exposing driver-specific information.
- Tuning kernel parameters.

# Creating a Procfs



## Steps to Create a Procfs Entry :

1. Use `proc_create()` to create an entry in `/proc`.
2. Define read/write operations using `proc_read` and `proc_write` functions.

# R/W Operations



## Code Summary:

Create /proc/simple\_proc\_entry.

### •Read Operation:

Use cat to read the value.

```
cat /proc/simple_proc_entry.
```

### •Write Operation:

Use echo to write a value.

```
echo 123 > /proc/simple_proc_entry.
```

# Procfs Subdirectory



## Steps to Create a Subdirectory in Procfs:

1. Use `proc_mkdir()` to create a directory.
2. Use `proc_create()` to create a file inside the directory.

# Use Cases of Procfs



- **Monitoring Process Information:**
  - Exposing CPU usage, memory usage, etc.
- **Driver-Specific Information:**
  - Exposing driver states and statistics.
- **Tuning Kernel Parameters:**
  - Adjusting kernel behavior at runtime.

# Best Practices



- Use procfs for process-related or kernel information.
- Avoid creating too many proc entries to prevent clutter.
- Ensure appropriate permissions to prevent unauthorized access.
- Validate user inputs to avoid kernel crashes.



# Introduction to Sysfs



- **Objective:**

- Understand the purpose of sysfs in Linux.
- Learn how to create read/write interfaces using sysfs.
- Explore practical examples for device drivers.

# What is Sysfs?



## **Definition:**

- Sysfs is a virtual file system in Linux that exposes kernel objects (devices, drivers, etc.) to user space.

## **Mounted Location:**

- Usually mounted at /sys.

## **Purpose:**

- Provides a structured way to expose kernel attributes.
- Allows user-space programs to interact with kernel objects.

# Why Use Sysfs?



- **Advantages:**

- Human-readable interface.
- Organized hierarchy.
- Easily customizable.

- **Common Use Cases:**

- Exposing device attributes.
- Interacting with hardware parameters.
- Controlling kernel modules.

# Creating a Sysfs



## Steps to Create a Sysfs Entry:

1. Use `kobject_create_and_add()` to create a directory in `/sys`.
2. Use `sysfs_create_file()` to create attributes.
3. Define read/write operations using `show` and `store` functions.

# R/W Operations



## Code Summary:

Create `/sys/kernel/simple_sysfs/my_value`.

## Read Operation:

Use `cat` to read the value.

```
cat /sys/kernel/simple_sysfs/my_value.
```

## Write Operation:

Use `echo` to write a value.

```
echo 42 > /sys/kernel/simple_sysfs/my_value.
```

# Practical Use Cases



- **Exposing Device Attributes:**
  - Device temperature, voltage levels, etc.
- **Hardware Control:**
  - Enable/disable hardware features.
- **Driver Debugging:**
  - Exposing internal driver states for debugging.

# Best Practices



- Use sysfs for hardware-related attributes.
- Organize entries in a hierarchical structure.
- Ensure appropriate permissions to prevent unauthorized access.
- Validate user inputs to avoid kernel crashes.

Feedback Time





Thank You