

Kernel Debugging

Team Emertxe



Introduction to Kernel Debugging



- Kernel debugging is critical for diagnosing issues in Linux kernel code.
- It involves identifying and fixing errors in:
 - Device drivers
 - System calls
 - Kernel modules

Challenges:

- No direct access to kernel memory
- Crashes can make systems unresponsive
- Requires specialized tools and methods

Debugging Methods



Debugging Methods Overview



1. Print-Based Debugging (using `printk`)
2. Crash Dump Analysis (using tools like `kdump`, `crash`)
3. Dynamic Probes (`kprobes`, `Uprobes`, `Ftrace`)
4. Interactive Debugging (using `KGDB`, `GDB`)
5. Static Code Analysis (tools like `sparse`, `smatch`)
6. Real-Time Debugging (using `perf`, `trace-cmd`)
7. Address2line Method

Kernel Logs and Printk



- `printk()` is the most common method for kernel debugging.
- Logs are available in `/var/log/messages` or `dmesg`.

Advantages:

- Simple to use
- Works in most environments

Disadvantages:

Requires code modifications

Not suitable for time-sensitive bugs

Example:

```
printk(KERN_INFO "Driver loaded successfully\n");
```

Crash Dump Analysis



- Used to analyze the state of the kernel when it crashes.
- Tools:
 - **kdump**: Captures kernel memory during a crash
 - **crash**: Analyzes crash dumps

Steps to Enable kdump:

1. Install kdump package
2. Configure `/etc/kdump.conf`
3. Enable and start kdump service

Example Command:

```
crash /var/crash/vmcore /usr/lib/debug/vmlinux
```

Dynamic Probes



Kprobes:

- Allows you to insert probes into running kernel code.
- Used to gather diagnostic data without recompiling the kernel.

Up probes:

- Similar to Kprobes but for user-space applications.

Ftrace:

- Provides function tracing within the kernel.

KGDB (Kernel GNU Debugger)



- KGDB allows you to debug the kernel using GDB.
- Requires two machines:
 - **Host:** Runs GDB
 - **Target:** Runs the kernel being debugged

Setup Steps:

1. Enable KGDB in kernel configuration.
2. Connect host and target via serial or network.
3. Start GDB on the host machine.

Static Analysis Tools



Sparse: Identifies common programming errors in kernel code.

Smatch: Static analysis tool tailored for the Linux kernel.

Advantages:

1. Finds issues without running the code.
2. Identifies potential security vulnerabilities.

Real-Time Kernel Debugging



- Tools like [perf](#), [trace-cmd](#), and [systemtap](#) enable real-time debugging.
- Useful for performance analysis and monitoring live systems.

Address2line Method



- address2line converts addresses from a crash dump or dmesg output to the corresponding source file and line number.
- Helps map kernel panic or OOPS addresses to specific code locations.

Steps to Use address2line:

- Identify the address from dmesg or crash dump.
- Use address2line with the kernel image:

Example Command:

```
address2line -e /usr/lib/debug/vmlinux-$(uname -r) 0xffffffff810c33a0
```

Advantages:

- Quickly pinpoints the source of the issue.
- No need to modify the kernel code.

Kernel Debugging Using QEMU and GDB



Introduction to GDB



The **GNU Debugger (GDB)** is a powerful tool used to debug programs by allowing you to:

- Set breakpoints
- Inspect variables
- Step through code
- Analyze crashes

For **kernel debugging**, GDB can connect to a running QEMU virtual machine to debug the kernel code in real-time.

GDB for Kernel Debugging



When using GDB for kernel debugging:

- You need to run the Linux kernel inside QEMU with debugging enabled.
- GDB connects to QEMU using a remote protocol ([gdbserver](#)).
- You can set breakpoints and tracepoints inside the kernel code and custom drivers.

Common GDB Commands



- `gdb <binary>`: Start GDB with a specified binary.
- `target remote :1234`: Connect GDB to a remote target (QEMU in this case).
- `break <function>`: Set a breakpoint at a function.
- `continue`: Resume program execution.
- `next`: Step to the next line of code.

What is QEMU?



QEMU is an open-source emulator and virtualizer that allows you to:

- Run virtual machines on your host system.
- Emulate different architectures (x86, ARM, etc.).
- Connect GDB to a running virtual machine for debugging purposes.

Why use QEMU for kernel debugging?



- No need for physical hardware.
- Easy to set up and configure.
- Supports different architectures (x86, ARM, etc.).

Preparing Your Host System



```
sudo apt update
```

```
sudo apt install build-essential libncurses-dev  
bison flex libssl-dev libelf-dev qemu qemu-  
system gdb
```

```
sudo apt install qemu qemu-kvm qemu-system-x86  
qemu-system-x86_64 --version
```

Linux Kernel - Setup



```
mkdir ~/kernel_debug
```

```
cd ~/kernel_debug
```

```
wget
```

```
https://cdn.kernel.org/pub/linux/kernel/v5.x/linux-5.15.tar.xz
```

```
tar -xf linux-5.15.tar.xz
```

```
cd linux-5.15
```

Configure the Kernel for QEMU



make x86_64_defconfig

make menuconfig

Enable debugging options:

- Kernel hacking → Compile the kernel with debug info
- Kernel hacking → Enable kgdb
- Kernel hacking → Magic SysRq key

Save and exit.

Compile the Kernel

make -j\$(nproc)

Transfer the Kernel to QEMU



The compiled kernel image will be located at:
`arch/x86/boot/bzImage`

Transfer the Kernel to QEMU

The bzImage file needs to be transferred to the directory from where QEMU will be launched. For simplicity, copy it to your `~/kernel_debug` folder:

```
cp arch/x86/boot/bzImage ~/kernel_debug/
```

Root Filesystem (Rootfs) for QEMU



A root filesystem is required for the Linux kernel to boot properly. You can use a prebuilt root filesystem or create your own using buildroot.

Option 1: Download a Prebuilt Rootfs

Option 2: Create Your Own Rootfs Using Buildroot/Yocto

Option 3: Creating an Initramfs using mkinitramfs

Starting Linux with QEMU



Start QEMU

```
cd ~/kernel_debug
```

```
qemu-system-x86_64 -kernel bzImage -initrd  
initramfs.cpio.gz -append "console=ttyS0 root=/dev/ram  
rdinit=/bin/sh" -nographic -s -S
```

Explanation:

- nographic: Disables graphical output.
- s: Opens port 1234 for GDB.
- S: Stops the CPU at startup (waiting for GDB).
- initrd initramfs.cpio.gz: Specifies the initramfs image.

Connect GDB to QEMU



Open another terminal and run:

- `cd ~/kernel_debug`
- `gdb vmlinux`
- `target remote :1234`

You can pass the following arguments to QEMU for debugging:

```
-append "console=ttyS0 root=/dev/ram  
rdinit=/bin/bash"
```


Breakpoints and Tracepoints in GDB



To Set a Breakpoint:

- `break start_kernel`
- `continue`

To Inspect Variables:

- `print <variable_name>`

Using Yocto - runqemu



What is QEMU?

- QEMU is an emulator that lets you run a virtual machine with a different architecture from your host system.
- Yocto integrates with QEMU to let you test your built images without real hardware.

What is runqemu?

runqemu is a Yocto tool that simplifies launching a virtual machine using QEMU with a Yocto-built Linux image.

Benefits of runqemu

- Fast and easy way to test your Yocto images.
- Automatically handles image and kernel selection.
- Supports network configurations and file sharing.

Basic runqemu Command



Syntax:

```
runqemu <image> <machine> <options>
```

Example:

```
runqemu core-image-minimal qemux86-64 nographic
```

Explanation:

- core-image-minimal: The image you built with Yocto.
- qemux86-64: The machine architecture.
- nographic: Run without a graphical interface.

Passing Extra Parameters to QEMU



Using qemuparms to Pass Custom QEMU Options

```
runqemu core-image-minimal qemux86-64 nographic  
qemuparms="-s -S"
```

Explanation:

- **-s**: Starts QEMU's GDB server on port 1234.
- **-S**: Stops the virtual CPU at startup, waiting for GDB to connect.